

Cliente Exemplo S.A.

AI Transformation Roadmap

Part 3: Application Platform


ASSESSMENT	AI Maturity Assessment 2026 H1
INDUSTRY	Financial Services / Insurance
ASSESSMENT DATE	2026-05-08
FRAMEWORK VERSION	1.0.0
GENERATED ON	2026-05-08
SCOPE	Engineering Organization

Table of Contents

1. Executive Summary	..
2. Transformation Roadmap	..
2.1 Horizon 1: Foundation (Months 0-6)	..
2.2 Horizon 2: Scale (Months 6-18)	..
2.3 Horizon 3: Transform (Months 18-36)	..
3. Resources	..
4. Success Metrics	..
5. Risks & Mitigations	..
6. Recommended Next Steps	..

1. Executive Summary

This document presents the transformation roadmap for Application Platform as part of Cliente Exemplo S.A.'s AI-Native Software Delivery Maturity transformation journey. The roadmap outlines the strategic initiatives, effort, and timeline required to evolve from the current pillar maturity (L1.9) to the target level (L4.0) over a 36-month period.



1.1 Current State Assessment

Metric	Current Value	Target Value	Gap
Platform Adoption	40%	>85%	+45pp
MTTR	4-8 hours	<30 min	Significant
Uptime	99.5%	99.95%	+0.45pp

1.2 Pillar 3 Capability Scores

#	CAPABILITY	CURRENT	TARGET	GAP	PRIORITY
3.1	PE Cloud-Native Architecture	L2.6	L3.0	0.4	LOW
3.2	PE API Management	L0.0	L3.0	0.0	LOW
3.3	AI Application Development	L2.2	L3.0	0.8	LOW
3.4	Data Platform & Lakehouse	L0.0	L3.0	0.0	LOW
3.5	Agentic Applications	L0.8	L4.0	3.2	CRITICAL
3.6	PE Identity & Access Management	L0.0	L3.0	0.0	LOW
3.7	Multi-Cloud & Portability	L0.0	L3.0	0.0	LOW
3.8	Performance & Scalability	L0.0	L3.0	0.0	LOW
3.9	FinOps & Cost Optimization	L0.0	L3.0	0.0	LOW
3.10	PE Platform Adoption & Engineering Culture	L3.0	L4.0	1.0	LOW

2. Transformation Roadmap

The transformation unfolds across three time horizons — H1 (Foundation, Months 0-6), H2 (Scale, Months 6-18), and H3 (Transform, Months 18-36) — enabling progressive capability building while delivering incremental value at each stage.

2.1 Horizon 1: Foundation (Months 0-6)

Target: L1.9 → L2.5 · Establishing AI-Assisted Development Foundation

PE 3.1 Cloud-Native Architecture

Current: L2.6 → H1 Target: **L2.5** · Gap: -0.1

CURRENT STATE EVIDENCE

- 70% containerized
- Service mesh pilot

H1 INITIATIVES

- ✓ Plan modernization for legacy monoliths
- ✓ Roll out service mesh org-wide
- ✓ Adopt event-driven architecture for new services

SUCCESS METRICS

METRIC	CURRENT	H1 TARGET
Containerization	70%	90%

PE 3.2 API Management

Current: L0.0 → H1 Target: **L1.5** · Gap: 1.5

CURRENT STATE EVIDENCE

- No IDP
- Service catalog planned
- 1 team experimenting

H1 INITIATIVES

- ✓ Define IDP vision and target architecture
- ✓ Deploy minimum-viable service catalog
- ✓ Identify and onboard 3 pilot teams to IDP

SUCCESS METRICS

METRIC	CURRENT	H1 TARGET
Teams on IDP	0	3

3.3 AI Application Development

Current: L2.2 → H1 Target: **L3.0** · Gap: 0.8

CURRENT STATE EVIDENCE

- K8s production
- GitOps
- 90% auto-scaling

H1 INITIATIVES

- ✓ Pilot AI-driven workload placement
- ✓ Optimize resource utilization with predictive scaling
- ✓ Standardize multi-cluster patterns

SUCCESS METRICS

METRIC	CURRENT	H1 TARGET
Cluster Utilization	60%	75%

3.4 Data Platform & Lakehouse

Current: L0.0 → H1 Target: **L2.5** · Gap: 2.5

CURRENT STATE EVIDENCE

- 70% OpenAPI
- No contract testing

H1 INITIATIVES

- ✓ Mandate OpenAPI specs for 100% of services
- ✓ Deploy contract testing in CI
- ✓ Enable AI-driven API discovery

SUCCESS METRICS

METRIC	CURRENT	H1 TARGET
OpenAPI Coverage	70%	100%

3.5 Agentic Applications

Current: L0.8 → H1 Target: **L2.5** · Gap: 1.7

CURRENT STATE EVIDENCE

- Lakehouse deployed
- Partial DQ monitoring

H1 INITIATIVES

- ✓ Expand data quality monitoring to all critical pipelines
- ✓ Roll out data mesh principles to 3 domains
- ✓ Standardize data contracts

SUCCESS METRICS

METRIC	CURRENT	H1 TARGET
DQ Coverage	50%	85%

PE 3.6 Identity & Access Management

Current: L0.0 → H1 Target: **L3.0** · Gap: 3.0

CURRENT STATE EVIDENCE

- 5 production use cases
- A/B testing
- Feature store pilot

H1 INITIATIVES

- ✓ GA the feature store for all teams
- ✓ Expand MLOps to 15+ use cases
- ✓ Deploy model monitoring with drift detection

SUCCESS METRICS

METRIC	CURRENT	H1 TARGET
Production Models	5	15

3.7 Multi-Cloud & Portability

Current: L0.0 → H1 Target: **L3.0** · Gap: 3.0

CURRENT STATE EVIDENCE

- 100% IaC
- Module testing
- Drift detection

H1 INITIATIVES

- ✓ Pilot AI-assisted module generation
- ✓ Deploy policy-as-code with AI suggestions
- ✓ Standardize module versioning

SUCCESS METRICS

METRIC	CURRENT	H1 TARGET
Module Reuse	60%	75%

3.8 Performance & Scalability

Current: L0.0 → H1 Target: **L2.5** · Gap: 2.5

CURRENT STATE EVIDENCE

- Tags enforced
- Manual rightsizing

H1 INITIATIVES

- ✓ Deploy automated rightsizing
- ✓ Build predictive cost analytics
- ✓ Establish FinOps as a formal practice

SUCCESS METRICS

METRIC	CURRENT	H1 TARGET
Cost Variance	±15%	±5%

3.9 FinOps & Cost Optimization

Current: L0.0 → H1 Target: **L2.0** · Gap: 2.0

CURRENT STATE EVIDENCE

- Multi-AZ tier-1
- 20% multi-region
- Inconsistent DR

H1 INITIATIVES

- ✓ Standardize quarterly recovery testing
- ✓ Expand multi-region to 60% of services
- ✓ Deploy reliability scorecards

SUCCESS METRICS

METRIC	CURRENT	H1 TARGET
Uptime (P99)	99.5%	99.9%

PE 3.10 Platform Adoption & Engineering Culture

Current: L3.0 → H1 Target: **L3.0** · Gap: 0.0

CURRENT STATE EVIDENCE

- Strong CoP
- Product mindset
- Quarterly surveys

H1 INITIATIVES

- ✓ Deploy AI-driven platform feedback analytics
- ✓ Establish platform OKRs with AI baselines
- ✓ Run platform innovation sprints

SUCCESS METRICS

METRIC	CURRENT	H1 TARGET
Platform NPS	+30	+45

2.2 Horizon 2: Scale (Months 6-18)

Target: L2.5 → L3.0 · Scaling Platform Engineering & Agentic Workflows

Key Initiatives

IDP GENERAL AVAILABILITY

- ✓ Migrate all teams to IDP
- ✓ Establish golden paths for top 10 use cases
- ✓ Deploy AI-enhanced developer self-service

AGENTIC WORKFLOWS

- ✓ Roll out Agent Mode org-wide
- ✓ Deploy AI-driven pipeline optimization
- ✓ Implement AI test selection and prioritization

H2 Capability Targets

CAPABILITY	H1 EXIT	H2 TARGET	KEY ENABLER
3.1 Cloud-Native Architecture	L2.5	L3.0	Service mesh standardization
3.2 API Management	L1.5	L3.0	IDP general availability
3.3 AI Application Development	L3.0	L3.5	AI workload placement
3.4 Data Platform & Lakehouse	L2.5	L3.5	AI API design assistant
3.5 Agentic Applications	L2.5	L3.5	AI data quality
3.6 Identity & Access Management	L3.0	L3.5	Production drift management
3.7 Multi-Cloud & Portability	L3.0	L3.5	AI module generation
3.8 Performance & Scalability	L2.5	L3.5	AI cost optimization
3.9 FinOps & Cost Optimization	L2.0	L3.0	Cross-region resilience
3.10 Platform Adoption & Engineering Culture	L3.0	L3.5	AI platform analytics

2.3 Horizon 3: Transform (Months 18-36)

Target: L3.0 → L4.0 · Autonomous Operations and Self-Healing Systems

Key Initiatives

AUTONOMOUS SDLC

- ✓ Deploy multi-agent coding workflows
- ✓ Implement autonomous progressive delivery
- ✓ Achieve self-healing infrastructure

CONTINUOUS OPTIMIZATION

- ✓ AI-driven workload placement
- ✓ Predictive incident prevention
- ✓ Continuous compliance with AI verification

H3 Capability Targets

CAPABILITY	H2 EXIT	H3 TARGET
3.1 Cloud-Native Architecture	L3.0	L4.0 — Cloud-native by default with mesh-aware routing
3.2 API Management	L3.0	L4.0 — AI-powered platform with golden paths
3.3 AI Application Development	L3.5	L4.0 — Autonomous workload optimization
3.4 Data Platform & Lakehouse	L3.5	L4.0 — Autonomous API governance
3.5 Agentic Applications	L3.5	L4.0 — Self-serving data products
3.6 Identity & Access Management	L3.5	L4.0 — Self-improving model lifecycle
3.7 Multi-Cloud & Portability	L3.5	L4.0 — Autonomous infrastructure provisioning
3.8 Performance & Scalability	L3.5	L4.0 — Continuous autonomous optimization
3.9 FinOps & Cost Optimization	L3.0	L4.0 — Autonomous resilience orchestration
3.10 Platform Adoption & Engineering Culture	L3.5	L4.0 — Self-improving platform with AI insights

3. Resources

3.1 Technology Investments

The technologies below are recommended over the 36-month roadmap. Effort tier indicates the implementation cost of each — see legend below.

TECHNOLOGY	HORIZON	EFFORT	PURPOSE
Service Catalog & IDP MVP	H1	L	Self-service foundation
Service Mesh	H1	M	Multi-cluster routing
Feature Store GA	H1	M	MLOps maturity
IDP General Availability	H2	XL	Org-wide self-service
AI Workload Placement	H2	M	Predictive scaling
Autonomous Platform Operations	H3	L	Self-improving platform

3.2 Team & Skills

ROLE	FTES	HORIZON	FOCUS
Platform Engineering Lead	1	H1	IDP strategy & architecture
Platform Engineers	8	H1	IDP buildout
MLOps Engineers	3	H1	AI/ML platform maturation
Platform Engineers	12	H2	IDP GA & scale operations

3.3 Effort Summary by Horizon

HORIZON	DURATION	EFFORT	FOCUS AREAS
H1	0-6 months	L	IDP MVP, service mesh, MLOps GA
H2	6-18 months	XL	IDP GA, predictive operations
H3	18-36 months	L	Autonomous platform, AI insights
TOTAL	36 months	XL	Complete Application Platform transformation

EFFORT TIER LEGEND

S

Small
4-8 weeks · 2-3 FTE ·
single squad

M

Medium
2-4 months · 4-6 FTE · 1-2
squads

L

Large
4-9 months · 7-12 FTE · 2-3
squads

XL

Extra-Large
9+ months · 12+ FTE ·
cross-org

4. Success Metrics

4.1 Platform Engineering Metrics

METRIC	CURRENT	H1 TARGET	H2 TARGET	H3 TARGET
Teams on IDP	0	3	All	All + AI golden paths
Platform NPS	+30	+45	+55	>+60
Self-Service Time-to-Provision	Days	Hours	Minutes	Instant

4.2 AI Platform Metrics

METRIC	CURRENT	H1 TARGET	H2 TARGET	H3 TARGET
Production Models	5	15	40	>80
Model Drift Detection	Manual	Automated	Predictive	Self-correcting

4.3 Resilience Metrics

METRIC	CURRENT	H1 TARGET	H2 TARGET	H3 TARGET
Uptime (P99)	99.5%	99.9%	99.95%	99.99%
Multi-Region Coverage	20%	60%	85%	100%

5. Risks & Mitigations

RISK	IMPACT	PROBABILITY	MITIGATION
IDP adoption stalls	HIGH	MEDIUM	Pilot teams as internal champions, golden paths for top use cases, exec sponsorship
Legacy modernization complexity	HIGH	HIGH	Strangler fig pattern, dedicated modernization team, phased decommission
Cost overrun on platform investments	MEDIUM	MEDIUM	Quarterly FinOps reviews, effort tier gating, pilot before scale

6. Recommended Next Steps

6.1 Immediate Actions (Next 30 Days)

- 1 Define IDP target architecture and golden-path principles
- 2 Hire Platform Engineering Lead
- 3 Inventory existing platform capabilities and gaps
- 4 Identify 3 pilot teams for IDP early adoption
- 5 Procure service mesh tooling

6.2 H1 Launch Activities (Days 30-90)

- 1 Deploy minimum-viable service catalog
- 2 Onboard first 3 pilot teams to IDP MVP
- 3 Begin service mesh deployment to 30% of services
- 4 Launch feature store GA
- 5 Establish platform-as-product OKRs

6.3 Dependencies

- Part 1: Developer Productivity adoption metrics for IDP success measurement
- Part 2: DevOps Lifecycle observability investments shared with platform